

Database operation vs Transaction

Database operation → Insert, Update, delete

Transaction → Business operations

Funds transfer, New Account

Opening

= "A transaction is a set of DML operation that performs a business unit of work."

Ex: Transfer 50\$ from account A to B

1) Read(A)

2) $A = A - 50$

3) write(B)

4) Read(B)

5) $B = B + 50$

6) write(B)

DML - Data manipulation language

How transaction Begins and How it ends!



Transaction Begins = when we execute the first DML statement after opening the session.

Transaction End : Txn ends with COMMIT or ROLLBACK or DDL or JCL.

⇒ There are 2 type of txn.

i) Read only (only select)

ii) Read write (insert, update, delete and select are allowed)

Transaction Properties :-

A → Atomicity

C → Consistency

I → Isolation

D → Durability

(operation)

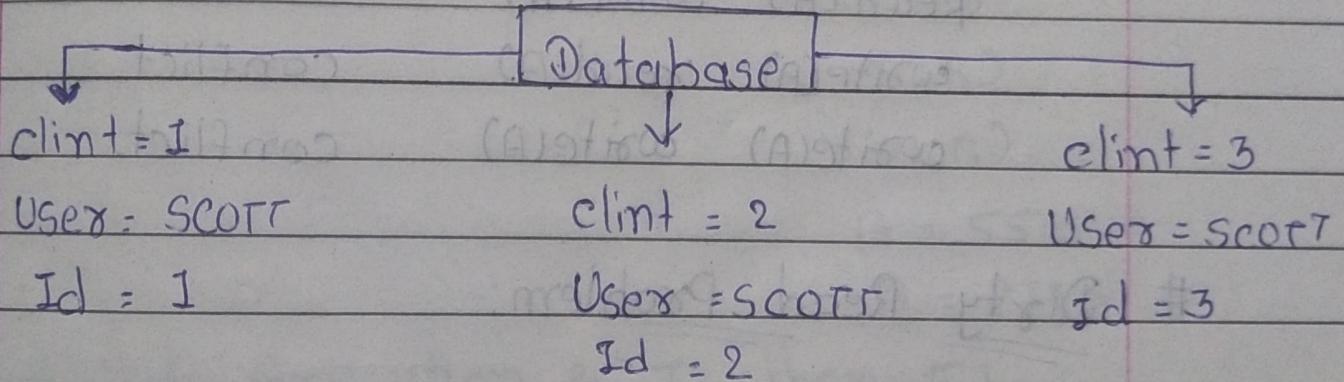
Atomicity :- Either all txns perform or not any txn perform.

Consistency :- when the transaction start and end the sum of money have to be equal.

Isolation :- Even T₁ and T₂ started concurrently the end result should be match at least one T₁ or T₂.

Durability : once transaction completed the data should be permanently saved.

Concurrency and Concurrency Control:



Insert into emp Insert into emp Insert into emp
 values (1,'a') values (2,'b') values (3,'c')

Transaction "T1" Transaction "T2" Transaction
 started Select * from emp "T3"
 Select * from emp;

1 row → (2,b) → (3,c)
 1 row (1,a)

Commit ;

Txn "T1" completed

Select * from emp;

2 rows (1,a)
 (2,b)
 select * from emp →
 (3,c)

→ Oracle uses Locking to control to concurrency. (for conflicting op.)

Two actions are said to Conflict if

- 1.) They belong to diff. transaction
- 2.) Both op. refers to the same db obj.
- 3.) one of them is write.

Ex:

T1 T2 CONFLICT or NOT

Read(A) Read(A)

Normal

Read(A) Write(A)

Conflict

Write(A) Read(A)

Conflict

Write(A) Write(A)

Conflict

Dirty Read Problem:-

Ex: T1 → Transfers 500 \$ from A to B

T2 → Add 10% interest to A and B

initial value

$$A = 1000$$

$$B = 2000$$

Serial Execution of T1 and then T2

T1

 $A = 1000$ $B = 2000$

R(A)

W(A)

R(B)

W(B)

 $A = 500$ $B = 2500$

T2

R(A)

W(A)

R(B)

W(B)

 $A = 550$ $B = 2750$

Serial Execution of T2 then T1

T1

T2

 $A = 1000$ $B = 2000$

R(A)

W(A)

R(B)

W(B)

 $A = 1100$ $B = 2200$

: begin

: end

R(A)

W(A)

R(B)

W(B)

 $A = 600$ $B = 2700$

Concurrent Execution :- T1 & T2

T1

 $A = 1000$ $B = 2000$

R(A)

W(A)

 $A = 500$ $B = 2000$

02FS = 8

T2

00RC = 9

R(A)

W(A)

R(B)

W(B)

 $A = 550$ $B = 2200$

R(B)

W(B)

 $A = 550$ $B = 2700$

→ Concurrent execution is not a Serializable schedule.

Unrepeatable Read [RW conflicts]

T1 = Reads 'A' two times & assigns to B cmd c

T2 = Increment two times value of A

A = 5

Serial Execution

T1 → T2

T2 → T1

T1

T2

A = 5

R(A)

W(A)

R(A)

W(C)

A = 5

B = 5

C = 5

W(A)

W(A)

A = 7

B = 7

C = 7

T1

T2

A = 5

W(A)

W(A)

A = 7

R(A)

W(B)

R(A)

W(C)

A = 7

B = 7

C = 7

Concurrent Execution

T1 & T2

T1

T2

R(A)

W(B)

 $B = 5$

commit W(A)

W(A)

 $A = 7$

R(A)

W(C)

 $C = 7$

	$A = 7$	
ST \rightarrow ST	$B = 5$	ST \leftarrow IT
ST	$C \neq 7$	IT

$\Rightarrow \Leftarrow$ same Acc. Value must be same if it differ then it called unrepeatable.

Blind Write Problem (Last update):

Ex: Akbar [A] and Birbal [B] are 2 employees and Trans. T1 & T2 wants to equal their Salaries.

$T_1 \rightarrow$ Equals A & B salary to 1111 ₹

$T_2 \rightarrow$ Equals A & B salary to 2222 ₹

$$A = 0$$

$$B = 0$$

serial Execution :-

$T_1 \rightarrow T_2$

$$A = 0$$

$$B = 0$$

$$R(A)$$

$$W(A)$$

$$R(B)$$

$$W(B)$$

$$A = 1111$$

$$B = 1111$$

$$R(A)$$

$$W(A)$$

$$R(B)$$

$$W(B)$$

$$A = 2222$$

$$B = 2222$$

$T_2 \rightarrow T_1$

$$A = 0$$

$$B = 0$$

$$R(A)$$

$$W(A)$$

$$R(B)$$

$$W(B)$$

$$A = 2222$$

$$B = 2222$$

$$R(A)$$

$$W(A)$$

$$R(B)$$

$$W(B)$$

$$A = 1111$$

$$B = 1111$$

Concurrent :-

T1

T2

 $A = 0$ $B = 0$

R(A)

W(A)

 $A = 1111$ $B = 0$

R(B)

W(A)

R(B)

W(B)

 $A = 2222$ $B = 2222$

R(B)

W(B)

 $A = 2222$ $B = 1111$

\Rightarrow more than one updates occurred it lost the previous one update so it also called Lost update problem.

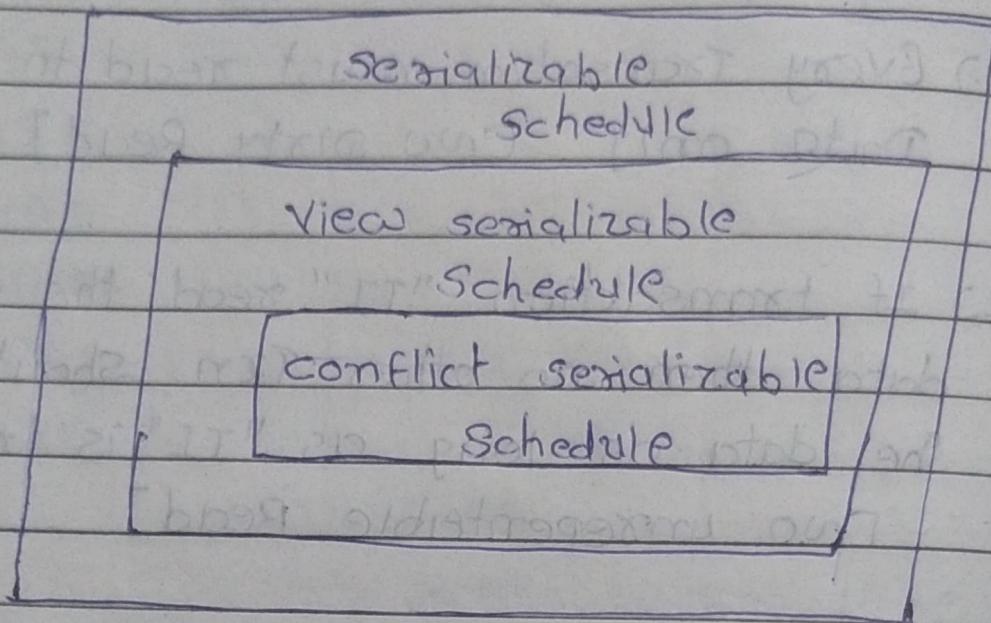
Guidelines to overcome all the 3 anomalies:-

- 1) Every Transaction must read to committed Data. only [No Dirty Read]
- 2) If transaction "T1" read the committed data then no other Trn. should modify the data as long as "T1" is in progress [No Unrepenentable Read]
- 3) if Trn "T1" has modified any data, then no other Trn is allowed to modify the data again as long as "T1" is in progress [No Blind Wait]

Testing for Serializability:-

- A schedule 's' is said to be serializable, if and only if the net effect is identical to some complete serial schedule over 's'
- As its difficult to verify whether the given schedule

All schedules :-



1.) Conflict Serializability :-

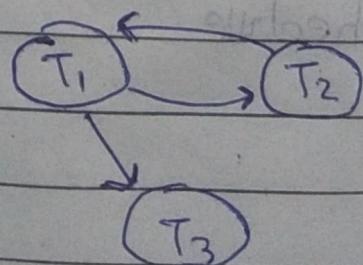
1.) draw a Precedence graph :-

(1.) A node for each committed transaction in S.

(2.) An arc from T_i to T_j if an action T_i Precedes and conflict with one of T_j 's actions

(3.) A schedule 'S' is conflict serializable , if and only if , the graph is acyclic.

T_1	T_2	T_3



if cycle occurs :- not conflict
else :- conflict

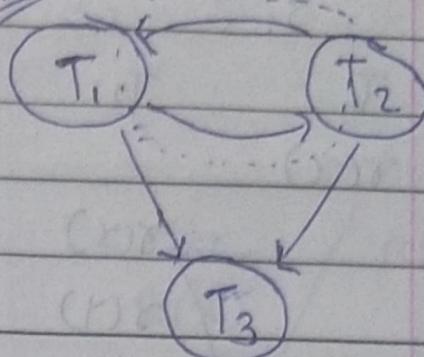
Page :

Date : / /

Ex:-

	T ₁	T ₂	T ₃
R(A)			
	W(A)		
	com.		
W(A)		v	
com.			
		W(R)	
		com.	

Conflict

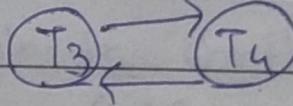


→ compare one Txn with every other Txn for finding conflict or not.

→ Cycle is formed ($T_1 \rightarrow T_2 \rightarrow T_3 \rightarrow T_1$) → not conflicting
Serializable

Ex:-

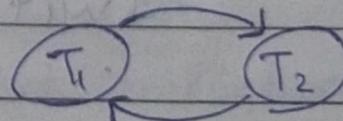
	T ₃	T ₄
R(Q)		
	W(Q)	
W(Q)		

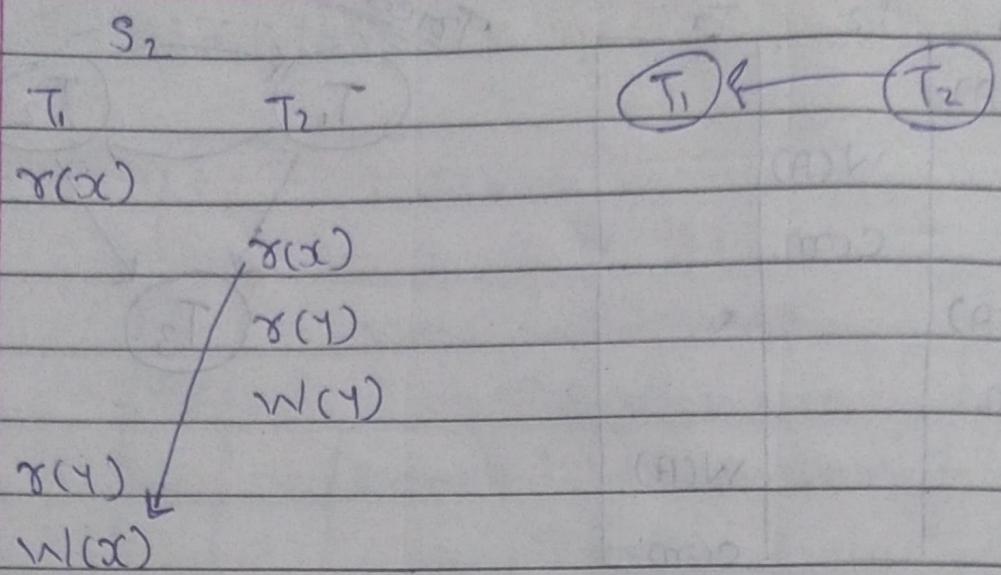


not-conflict serializable

Ex:-

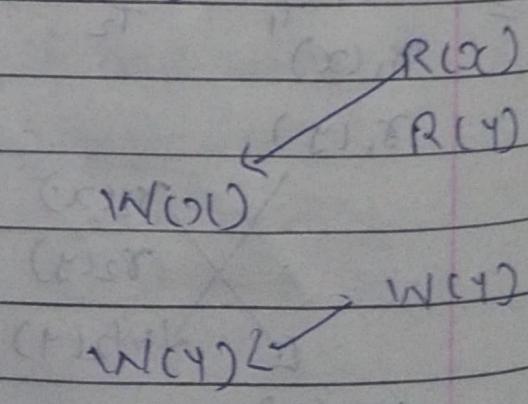
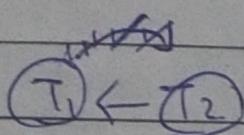
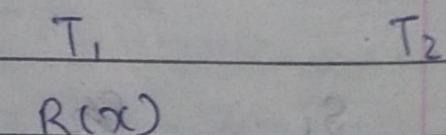
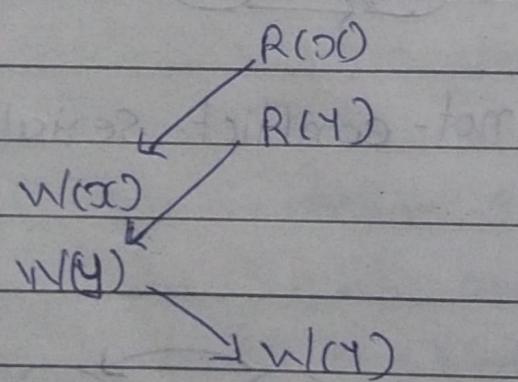
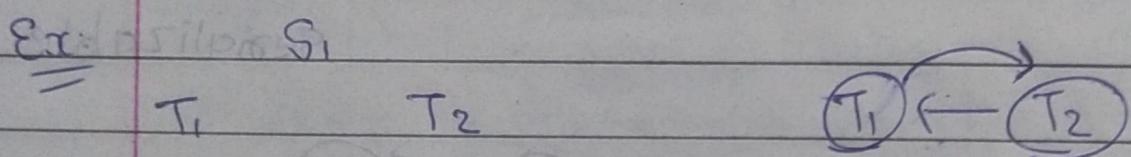
	S ₁	S ₂
$\gamma_1(x)$		
$\gamma_1(y)$		
	$\gamma_2(x)$	
	$\gamma_2(y)$	
		$w_2(y)$
	X	
$w_1(x)$		

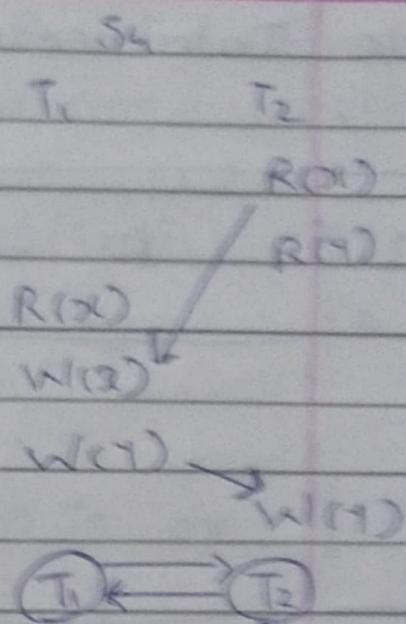
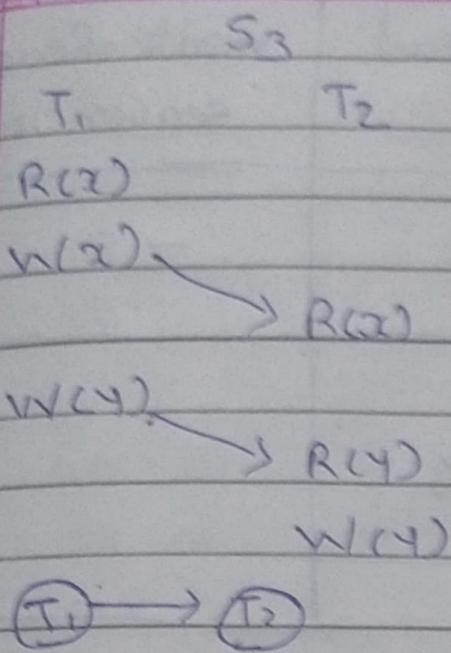




$S_1 \rightarrow$ not conflict serializable

$S_2 \rightarrow$ conflict serializable

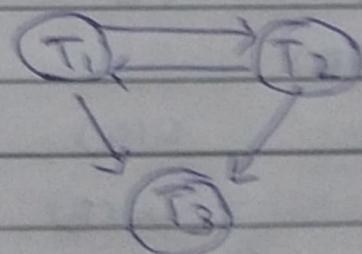
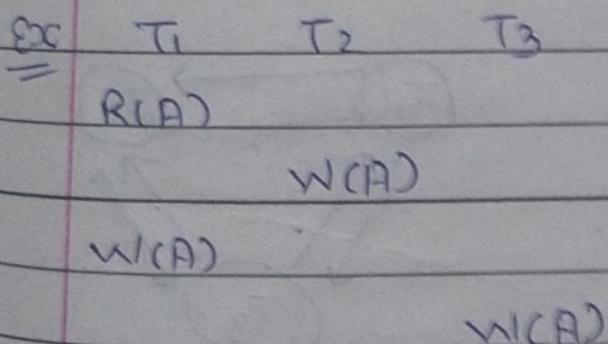




→ S_2 and S_3 conflict serializable

View Serializability :-

1.) whether the schedule is conflict serializable or not.



not-conflict serializable

→ for view serializable create table

SNO

Database obj

Initial Read

Write-Read
sequencing

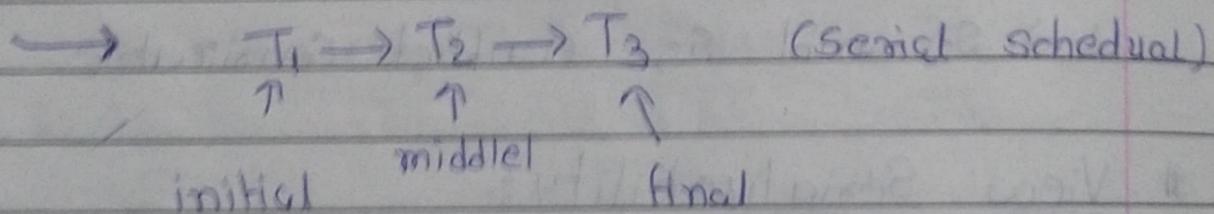
1

A

T₁

final

mark

T₁ → T₃T₃

\rightarrow view serializable above ex.

Ex-2T₁T₂T₃

R(A)

R(B)

W(B)

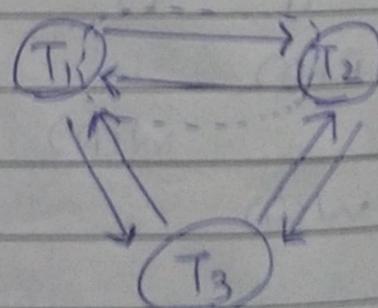
R(A)

W(A)

R(B)

R(A)

W(A)



(not-conflict
serializable)

SNO	Db. obj.	initial Read	W-R	Final
			sequencing	write

1.	A	T1	Yes ($T_3 \rightarrow T_2$)	T2
2.	B	T2	($T_2 \rightarrow T_1$, $T_2 \rightarrow T_3$)	T2

for obj. A :- $T_1 \rightarrow T_3 \rightarrow T_2$

for obj. B :- $T_2 \rightarrow T_1 \rightarrow T_3$ or

$T_2 \rightarrow T_3 \rightarrow T_1$

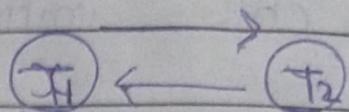
→ check order and match both object serial order.

→ if it is not matching then it is not view serializable.

→ if it is matching then it is view serializable.

Ex:3

T₁ T₂



R(x)

w(y)

R(y)

w(y)

R(y)

w(y)

R(x)

w(x)

(not - conflict)

1 x

2 y

T₁

T₂

T₁ → T₂

T₂ → T₁

T₂

T₁

x : T₁ → T₂

y : T₂ → T₁

(not view serial)

Serial : T₁ → T₂

T₁ T₂

x = 100

R(x)

y = 100

x = x - 10

w(x)

R(y)

y = y + 10

w(y)

R(y)

y = y + 20

w(y)

R(x)

x = x - 20

w(x)

Page No:

Date:

Page :

Date : / /

$$x = 90$$

$$y = 110$$

$$x = 80$$

$$y = 120$$

$$x = 70$$

$$y = 130$$

Serial order $T_2 \rightarrow T_1$

$$x = 80$$

$$y = 120$$

$$x = 70$$

$$y = 130$$

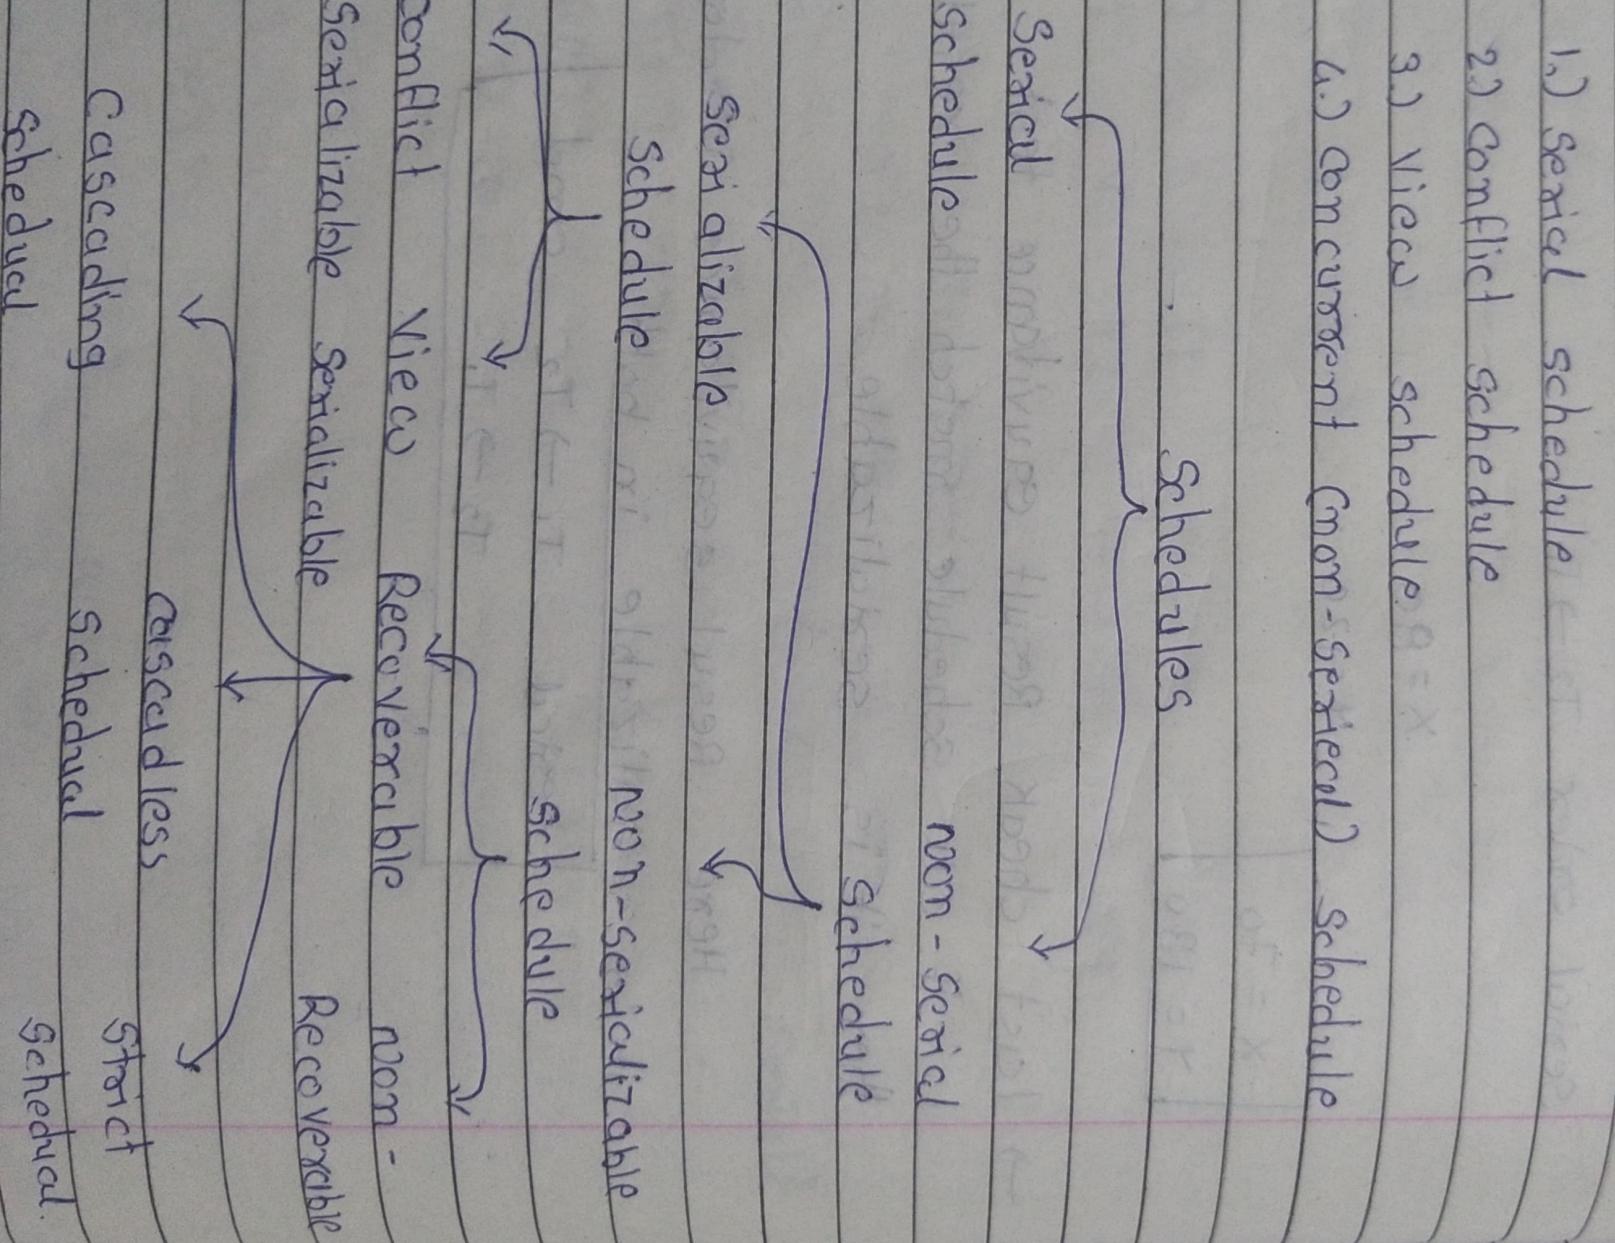
→ Last check Result equivalence test.
if all schedule match the result
then it is serializable.

Here Result equivalence is done and
it is serializable in both

serial $T_1 \rightarrow T_2$ and
 $T_2 \rightarrow T_1$

Types of Schedules :-

* Schedule :- "The order



non-serializable

non-serializable schedules :-

Two types :-

1.) Recoverable :-

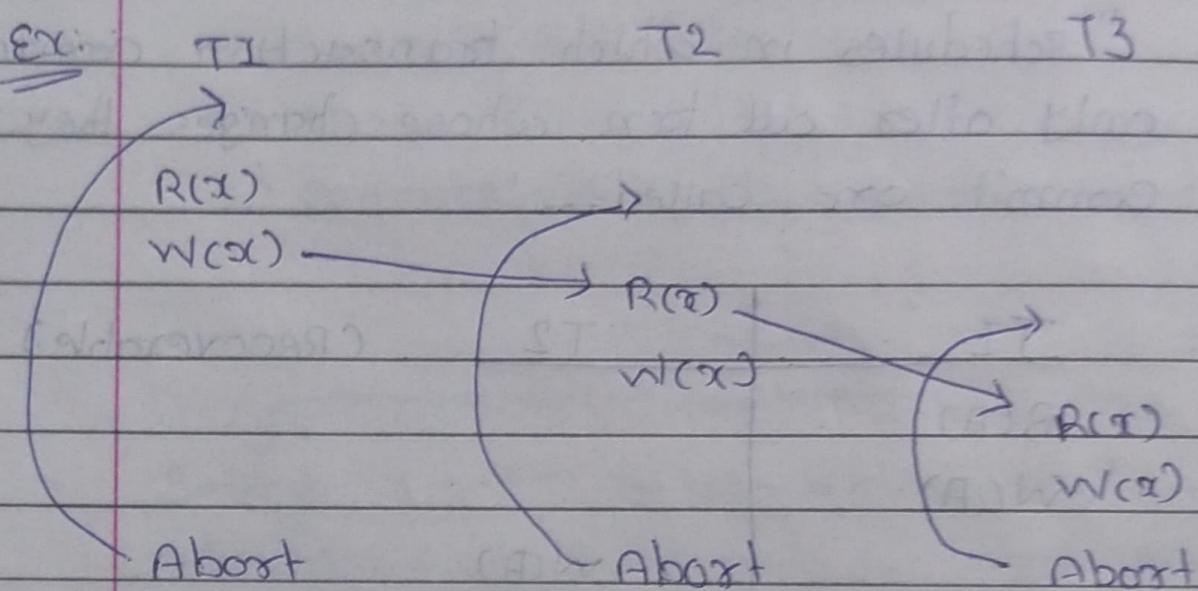
"schedules in which transaction commit only after all txn whose changes they read Commit are called..."

<u>Ex:</u>	<u>T1</u>	<u>T2</u>	<u>(Recoverable)</u>
	R(A)		
	W(A)		
		W(A) R(A)	
	commit		commit.

- If only txn changes the cmd it fails then and only then it's not Recoverable.
- If only Read op. perform and other txn changes and 1st txn fail also it called Recoverable.

Cascading schedule :- (Avoid Cascading aborts) (ACA)

"when there is a failure in one trn and this leads to the rolling back or aborting other dependent trn. it called ---"



Cascade less schedule :-

"~~if~~ schedules in which trn read values after all trn whose changes they are going to read commit are called."

\Rightarrow other (T2) trn can't read T1 trn (CR(A))

before the Trn T1 commit

Ex:

T1	T2	
	R(A)	→ is one R(A) will need &
	W(A)	other transaction will R(A) on which.
R(B)		
W(C)		
commit		
R(A)		
commit		

strict schedule :-

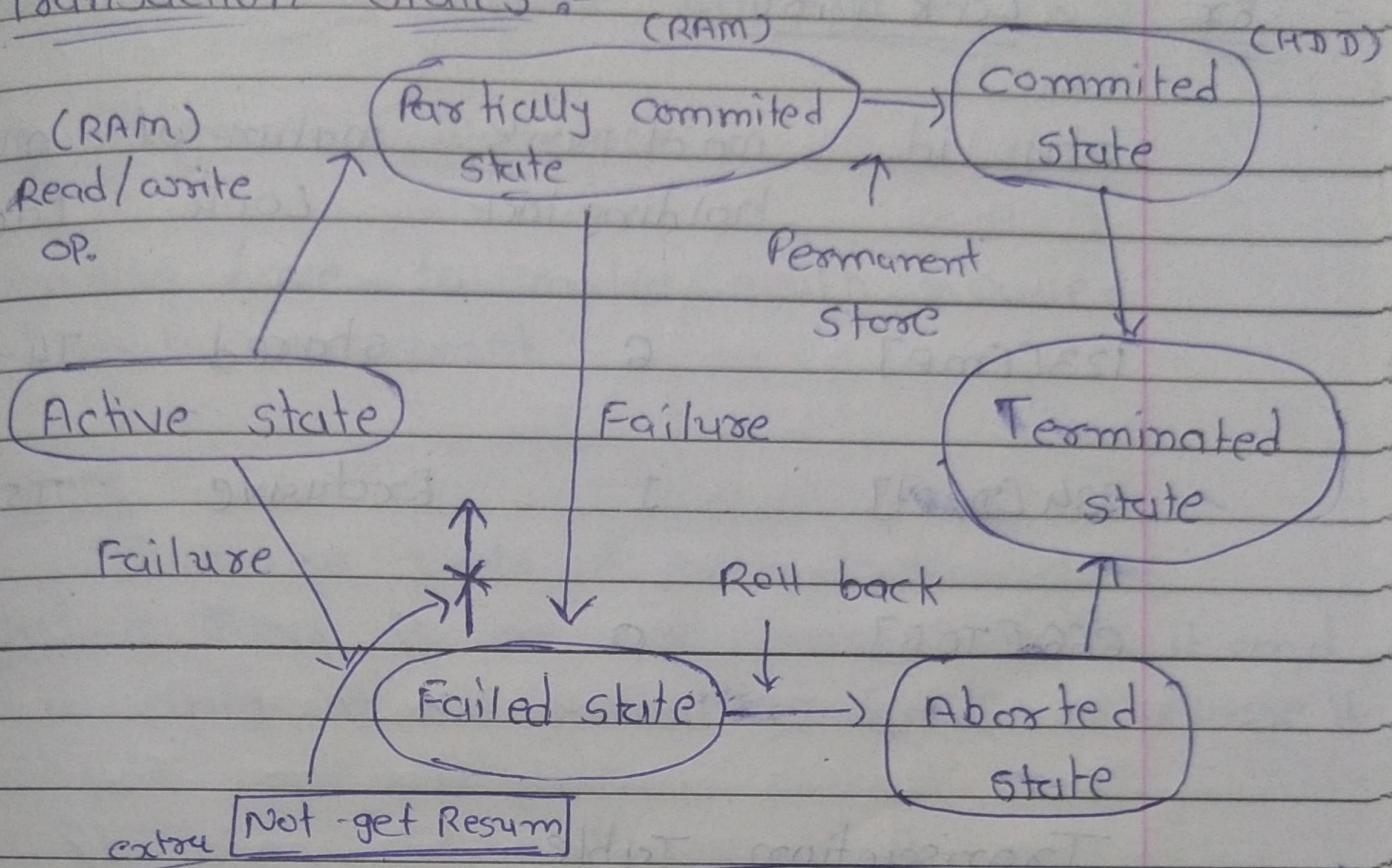
⇒ other transaction are not allowed R(A) or W(A) before the commit.

↪ if there is a change after that R(A) or W(A) not allowed. or value change not allowed until the transaction T1 commit.

Ex:

T1	T2	
R(A)	it will be not	→ is one value change
	R(A)	will need R(A) so
W(A)	conflict.	W(A) on which is
Commit,		commit such.
W(A)		
R(A)		
Commit		ToPa

Transaction States :-



[Txn always Restart]

[never Resum]

Intro. to Lock Management :-

↳ can be granted by Lock manager.

1.) Shared Lock. → Read allowed

2.) Exclusive Lock. → R,W allowed

↳ Some Txn T1 having shared lock than other Txn only allowed shared lock.

↳ any other case not allowed ~~Ejemplo~~ khori.

Ex:

Lock Table

obj. id	no of Trn's holding lock	nature of lock	queue of lock
123 [EMP]	6	shared	T4 → T5 → T6
234 [DEPT]	1	Exclusive	TS → T6
666 [JOB]	0		

Transaction Table

Txn_ID List of Locks

T1 R[EMP], W[DEPT]

T2 W[EMP] ← R[EMP]

T3 -

S = Shared lock

X = Exclusive lock

Page 1

Data 1 1 1

Deadlock & Specialize Locking :-

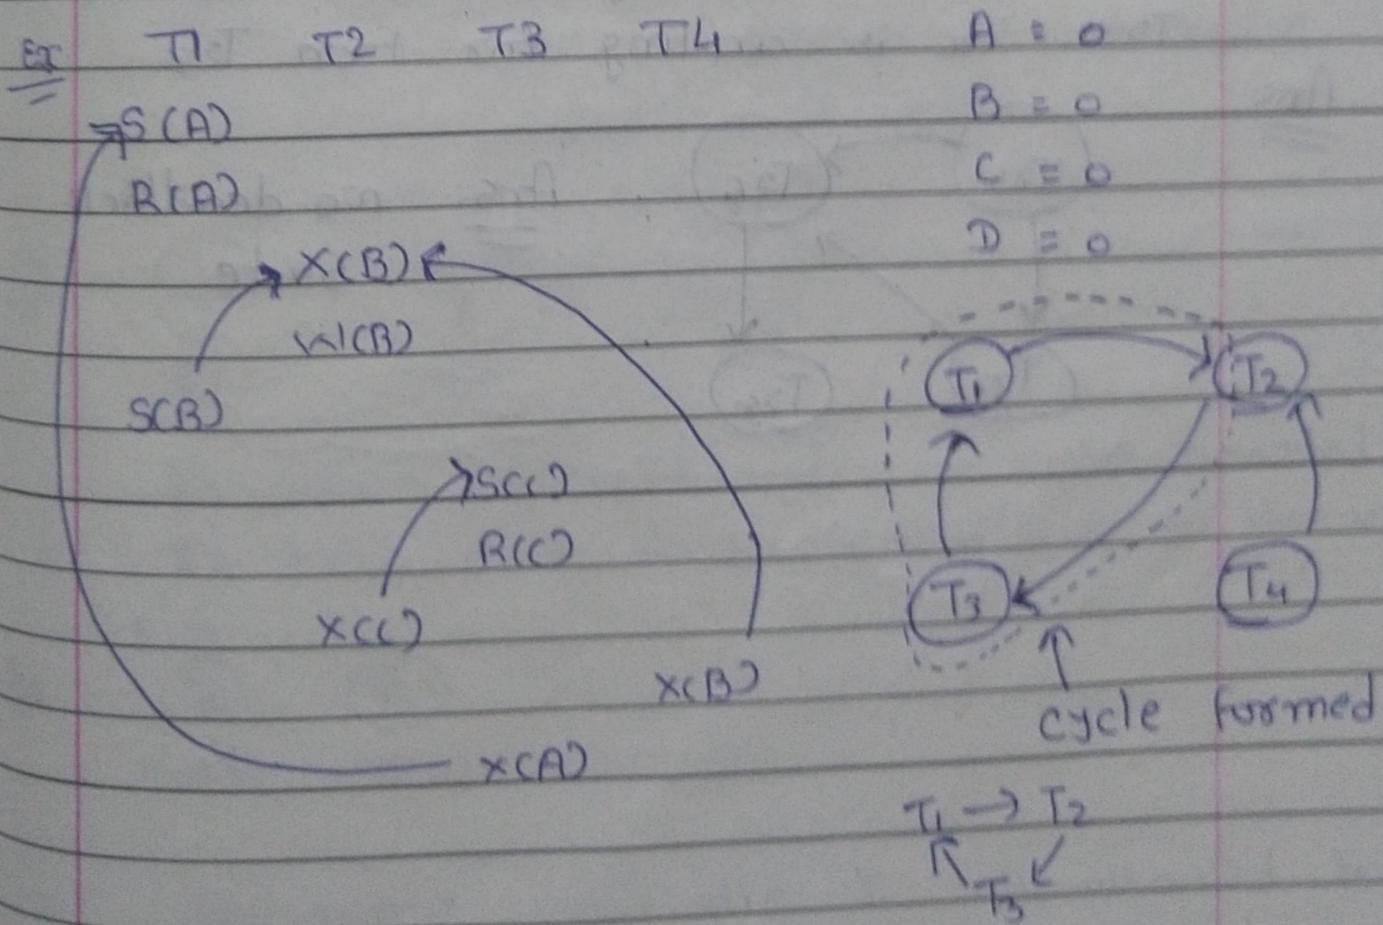
→ Declining with Deadlock:-

→ we have to make graph called "waits-for-graph"

Steps :-

→ make node of every txn.

imp → There is an arc from T_i to T_j , if and only if, T_i is waiting for T_j to release it.



Ques Que :-Ex:-T₃ T₄

X(B)

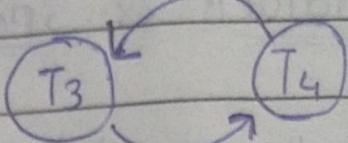
R(B)

B = B - 50

W(B)

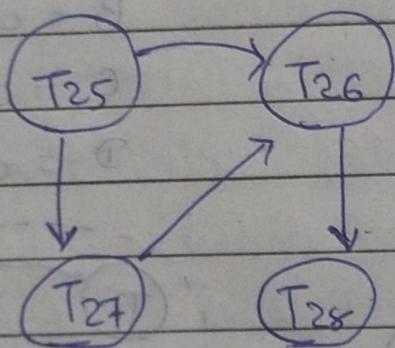
S(A)
R(A)
SC(B)

X(A)



Ans:- Result in Deadlock

- Ex:- T₂₅, T₂₅ is waiting for token T₂₆ and T₂₇
 T₂₆, T₂₇ is waiting for token T₂₆
 T₂₆, T₂₆ is waiting for token T₂₈

Ans:Ans: no deadlock
state.

DeadLock Prevention :-

→ we can prevent deadlock by giving each txn a Priority to each txn.

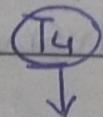
two methods available:-

1.) wait-die :- if T_i has higher priority it is allowed to wait, otherwise, it is aborted.

2.) round-wait :- If T_i has higher priority abort T_j ; otherwise, T_i waits.

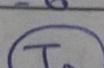
$P = 2$

$P = 3$



By using round-wait T_4 will be aborted
 T_3 txn. and first T_4 txn will go through.

$P = 7$ $P = 6$



$X(A)$ $S(A)$

By using wait-die T_1 will be allowed to wait and if its $P < 6$ it will be aborted.

Starvation :- Finite time waiting.

Page : / /

Date : / /

Shared - Exclusive locking

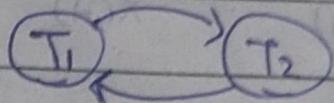
(1) Shared lock (S) :- S(A) :-

(2) Exclusive lock (X) :- X(A) :-

<u>Ex:</u>	<u>T1</u>	<u>T2</u>
	S(A)	
	B(A)	
<u>unlock</u> <u>The Lock</u>	$\rightarrow U(A)$	
		X(A)
		B(A)
		W(A)
		U(A)

→ Deadlock in shared-Exclusive locking:-

<u>Ex:</u>	<u>T1</u>	<u>T2</u>
	X(A)	
	X(B)	
	X(B)	X(A)



Irrecoverable schedule 3-

TI	T2
X(A)	new slot + slot agenda
R(A)	old slot
W(A)	old slot + slot agenda
U(A)	old slot
X(A)	
R(A)	OT
W(A)	OT
U(A)	OT
Commit	cancel

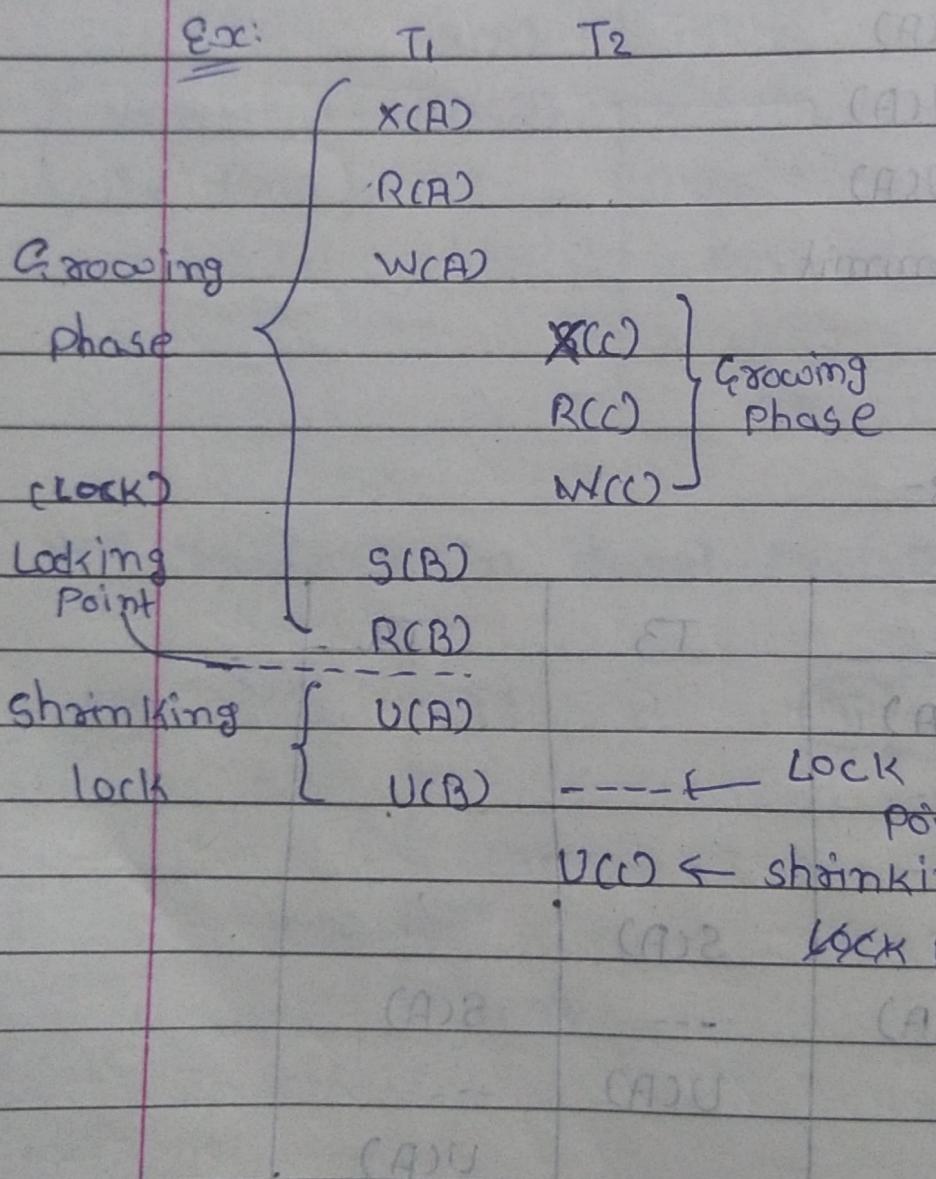
→ Starvation :-

There is starvation in Town A.

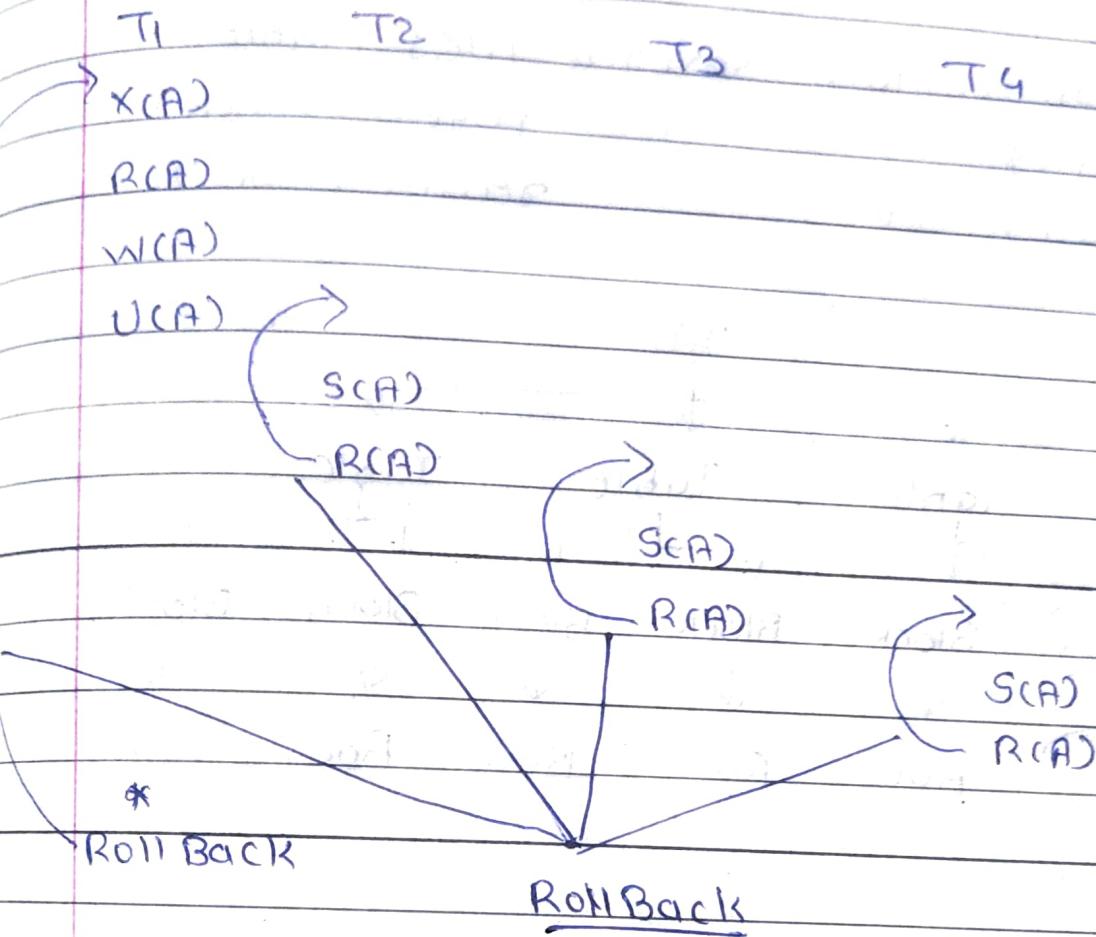
2-Phase locking :-

Growing lock : lock are acquired and no locks are released.

Shrinking lock : locks are released and no locks are acquired.



Cascading RollBack :-

# 1) Strict 2PL :-

→ Satisfy 2PL and X(exclusive) lock hold until commit/abort.

2) Rigorous 2PL :-

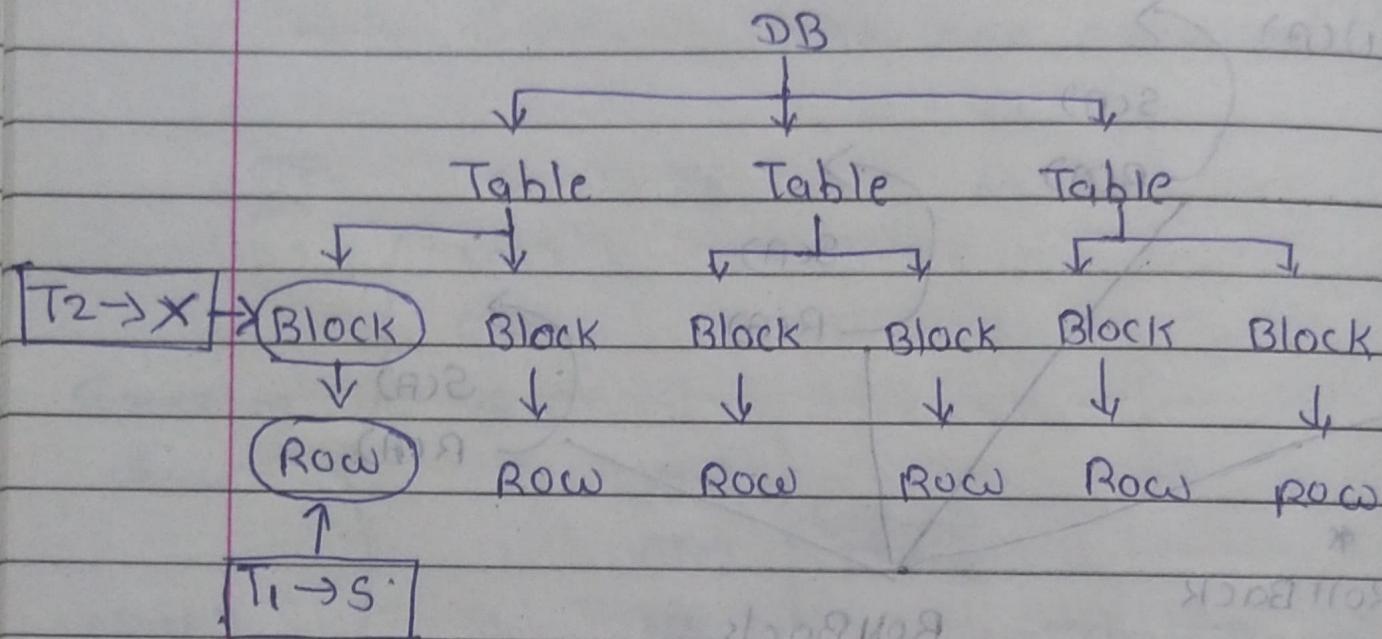
→ Satisfy 2PL and all S(shared) and X(exclusive) locks hold until commit/abort.

Multiple - Granularity Locking

→ DB contain several files (table).

→ files " " Pages (block)

→ Pages " " records (row)



→ Shared lock (S) } locks used to lock the
Exclusive lock (X) } object (Row)

→ Intention shared (IS) } locks used to lock
Intention Exclusive (IX) } ancestors (Table,
block)

→ IS conflict with X locks

→ IX conflict with S and X locks

- Lock must be acquired from Root-to-leaf
- Lock must be released from Leaf-to-Root

Log based recovery : (deferred database modification)

Syntax:

< Transaction no., object, new value >

T1

A = 100

R(A)

< T1, START >

A = A + 100

W(A)

< T1, A, 200 >

R(B)

< T1, B, 400 >

B = B + 200

< T1, COMMIT >

W(CB)

A = 200

commit

B = 400

* failed

Log based recovery : (Immediate database modification)

< T1, START >

< T1, A, 100, 200 >

< T1, B, 200, 400 >

< T1, COMMIT >

→ If transaction is not committed and it failed then we will do UNDO operation else we do REDO operation.

Timestamp Protocol

Before that some known :-

TS = Time stamp

$TS(T_i)$ = T_i 's Time stamp

→ Whenever a transaction come into system the unique value assigned to it.

RTS = Read time stamp (Last transaction TS)

WTS = write time stamp (Last transaction TS)

Time stamp Protocol Rules

1) R(A)

A) $WTS(A) > TS(T_i)$ then roll back T_i .

B) else execute $RTS(A) = \max\{RTS(A), TS(T_i)\}$

2) W(A)

A) $RTS(A) > TS(T_i)$ } Rollback T_i if and only if both true

B) $WTS(A) > TS(T_i)$