

## TECHNICAL RESOURCES

Connecting our research communities to information, tools, and resources across HPC, AI and Big Data.



### Intro to GPU Programming

- Video: Introduction to NVIDIA GPU Computing
- Presentation: Present and Future of Accelerated Computing Programming Approaches
- Video: ACM Winter School 2019 on HPC IT Kanpur

### Libraries

- Documentation: Math and communication libraries.
- Presentation: How CUDA Math Libraries Can Help You Unleash the Power of the New NVIDIA A100 GPU.
- Video: Recent Developments in NVIDIA Math Library
- Presentation: High-Performance Sparse Linear Algebra on NVIDIA GPUs with cuSPARSE.

### ARM

- Presentation: Port, Profile, and Tune HPC Applications for Arm-based Supercomputers.
- Presentation: Port, Profile, and Tune HPC Applications for Arm-based Supercomputers.
- Web Page: NVIDIA Arm HPC Developer Kit.

### Programming Models

#### OpenACC

- Web Page: OpenACC.org Resources
- Forum: OpenACC Slack Channel
- Presentation: Zero to GPU Hero with OpenACC
- Presentation: Directive-Based Programming with OpenACC
- Training Series: OpenACC Training Series
- Video: Directive Based GPU Programming
- GitHub: OpenACC Training Materials
- Docker Container: OpenACC Training Materials
- Online Course: Fundamentals of Accelerated Computing with OpenACC (Fee-based)
- Video: OpenACC Multi-GPU

#### Standard Languages

- Presentation: Standard Language Parallelism
- Presentation: Developing HPC Applications with Standard C++, Fortran, and Python
- Presentation: The NVIDIA C++ Standard Library
- Presentation: Shifting through the Gears of GPU Programming: Understanding Performance and Portability Trade-offs

#### CUDA

- Training Series: CUDA Training Series
- Presentation: CUDA on NVIDIA Ampers GPU Architecture: Taking Your Algorithms to the Next Level of Performance
- Presentation: CUDA: New Features And Beyond
- Presentation: Developing CUDA Kernels to Push Tensor Cores to the Absolute Limit on NVIDIA A100
- Online Course: Fundamentals of Accelerated Computing with CUDA C++ (Fee-based)
- Presentation: How CUDA Programming Works (ETC Fall '22), How CUDA Programming Works (ETC Spring '22)
- Presentation: CUDA Programming Model for Hopper Architecture
- Presentation: Optimizing CUDA Applications for NVIDIA Hopper Architecture
- Presentation: Enabling Hopper-Specific Optimizations in CUDA Applications
- Presentation: Performance Optimization with Modern CUDA Programming Techniques

#### CUDA Fortran

- Documentation: Programming Guide

#### Python

- Presentation: CuPy Overview: NumPy Syntax Computation with Advanced CUDA Features
- Presentation: Accelerating Python with CUDA
- Online Course: Fundamentals of Accelerated Computing with CUDA Python (Fee-based)
- Tutorial: Python Profiling
- Presentation: GPU Acceleration in Python
- Presentation: Legate: Scaling the Python Ecosystem
- Presentation: Scale Python and Numpy Performance with Legate
- Presentation: GPU Acceleration in Python using CuPy and Numba
- Presentation: Accelerate Computing with CUDA Python

#### Kokkos

- GitHub: Kokkos Repository
- GitHub: Kokkos Tutorials
- Forum: Kokkos Slack Channel
- Video: Kokkos: C++ Performance Portability for Production

#### RAJA

- GitHub: RAJA Repository
- Video: A Tutorial Introduction to RAJA

#### Multi-GPU

- Presentation: Multi-GPU Programming with MPI (a Magnum IO Session)
- Presentation: Multi-GPU Programming Models
- Presentation: A Partitioned Global Address Space Library for Large GPU Clusters
- Presentation: NVSHMEM: CUDA-Integrated Communication for NVIDIA GPUs (a Magnum IO session)
- Presentation: NVSHMEM: GPU-Integrated Communication for NVIDIA GPU Clusters
- Presentation: NCCL: High-Speed Inter-GPU Communication for LargeScale Training
- GitHub: Multi-GPU Programming Models

#### Tools

- Download: NVIDIA Nightly™ Systems, NVIDIA Night Compute, NVIDIA Night Graphics, NVIDIA Night Visual Studio Code Edition, NVIDIA Deep Learning Designer
- Documentation: Night Systems, Night Compute, Night Graphics, CUPTI, NVIDIA Tools Extension SDK (Nvtx), Compute Sanitizer, CUDA-memcheck, CUDA-GDB, Night Visual Studio Code Edition, Night Deep Learning Designer
- Video: NVIDIA Night Systems "Profiling GPU Applications with NVIDIA nSight Systems"
- Video: NVIDIA Night Compute "Introduction to NVIDIA Night Compute"
- Technical Session: CUDA is Evolving, and the Latest Developer Tools are Adapting to Keep Up
- Technical Session: It's Alive: CUDA in Visual Studio Code
- Technical Session: Optimizing CUDA Kernels in HPC Simulation and Visualization Codes Using NVIDIA Night Compute
- Technical Session: What the Profiler is Telling You: How to Get the Most Performance out of Your Hardware
- Technical Session: Performance Tuning CUDA Applications with the Roofline Model
- Technical Session: Roofline and NVIDIA Ampere GPU Architecture Analysis
- GitHub: Night Developer Tools Training
- Technical Blog: Analysis-Driven Optimization: Analyzing and Improving Performance with NVIDIA Night Compute
- Technical Blog: Custom Application Profile Timelines with NVTX
- Demo: NVIDIA Night Developer Tools
- Blog: Subverting CUDA More Efficiently with NVIDIA Compute Sanitizer
- Technical Session: How to Understand and Optimize Shared Memory Accesses using Night Compute
- Technical Session: Requests, Wavefronts, Sectors Metrics: Understanding and Optimizing Memory Bound Kernels with Night Compute

#### Data Science

- Web Page: RAPIDS overview
- Web Page: RAPIDS Getting Started (Hands-on labs and other materials)
- Forum: RAPIDS Community
- Presentations: Collection of RAPIDS Talks
- Online Course: Fundamentals of Accelerated Data Science with RAPIDS(Fee-Based)

### AI/Deep Learning Libraries, Frameworks, SDKs

#### Introduction to AI

- Presentation: Dive into Deep Learning
- Presentation: Do-It-Yourself Automatic Speech Recognition with NVIDIA Technologies
- Online Course: Fundamentals of Deep Learning for Computer Vision (Fee-Based)
- GitHub: Deep Learning Examples(The latest deep learning example networks for training and Inference.)

#### cuNumeric

- Presentation: cuNumeric and Legate: How to Create a Distributed GPU Accelerated Library
- Presentation: Evaluating Your Options for Accelerated Numerical Computing in Pure Python

#### CuDNN

- Documentation: CuDNN Developer Guide
- Presentation: cuDNN v8 New Advances in Deep Learning Acceleration: APIs, Optimizations, and How to Tackle the Future Challenges in Hardware and Software
- Presentation: Deep Learning Training with cuDNN

#### NVIDIA CLARA

- Container: NVIDIA Clara Train SDK
- Jupyter Notebooks: Intro to Clara Train SDK
- Jupyter Notebook: Clara Federated Learning
- Presentation: Clara Developer Day: Federated Learning Using Clara Train SDK
- Presentation: Clara Developer Day: Clara Train SDK Performance Walkthrough and Deep Dive
- Presentation: Clara Developer Day: Scalable and Modular Deployment Powered by Clara Deploy SDK
- Presentation: Clara Developer Day: Getting Started with Clara Train for High Performance & Iterative Experimentation with AutoML
- Web Page: NVIDIA Clara™ Parabricks Overview

#### DeepStream

- Presentation: Developing IVA Software Using NVIDIA DeepStream SDK
- Online Course: Getting Started with DeepStream for Video Analytics on Jetson Nano(Fee-Based)
- Online Course: AI Workflows for Intelligent Video Analytics with DeepStream(Fee-Based)

#### JAX

- Presentation: Seamlessly Scale Out Complex Compute Workloads with Multi-Node JAX
- Presentation: What's New in JAX

#### TensorRT

- Presentation: PyTorch-TensorRT: Accelerating Inference in PyTorch with TensorRT
- Presentation: Inference with TensorFlow 2 Integrated with TensorRT
- Presentation: TensorRT Inference with TensorFlow 2.0
- Online Course: Optimization and Deployment of TensorFlow Models with TensorRT(Fee-Based)

#### TensorFlow

- Presentation: What's New in TensorFlow-An Ecosystem Update
- Presentation: Extensions of TensorFlow-Based Computational Fluid Dynamics
- Online Course: Image Classification with TensorFlow: Radiomics - 1p19q Chromosome Status Classification(Fee-Based)
- Demo: Profiling TensorFlow using Tensorboard

#### Keras

- Presentation: Keras and TensorFlow: The Next Five Years
- Presentation: Porting Keras Lambda Layers to TensorFlow.js
- Online Course: Modeling Time Series Data with Recurrent Neural Networks in Keras(Fee-Based)

#### PyTorch

- Website: PyTorch
- Container: PyTorch
- Presentation: Latest Developments on PyTorch: Getting Started and Community Highlights
- Presentation: Dynamic Shapes First: Advanced GPU Fusion in PyTorch
- Presentation: Scaling Large Models Using PyTorch RPC
- Presentation: Named Tensors, Model Quantization, and the Latest PyTorch Features
- Demo: Profiling PyTorch with Night System

#### Containers

- Video: Making Containers Easier with HPC Container Maker.
- Presentation: Containers Democratize HPC.
- Online Course: High-Performance Computing with Containers(Fee-Based).
- Presentation: Simplifying and Accelerating HPC Workflows with NGC.

### Advanced

- Presentation: Inside the NVIDIA Ampere Architecture.
- Presentation: Optimizing Applications for NVIDIA Ampere GPU Architecture.
- Blog: NVIDIA Hopper Architecture In-Depth.
- Presentation: Inside the NVIDIA Hopper Architecture.

#### Navigate

- Events
- Attendees
- Mentors
- Resources
- Log In

#### About Open Hackathons

- Blog
- News
- Join Our Community
- Contact Us
- Terms and Conditions
- About

#### Company

- About OpenACC Organization
- Code of Conduct
- Privacy Policy
- Contact Site Administrator

#### Stay in Touch



